CN7023 - Artificial Intelligence & Machine Vision

Project Topic: Rice Variety Classification Using MobileNetV2 and InceptionV3

Name: Vinitha Abhinashree M

Email: vinitha.shree9923@gmail.com

Dataset Description

Dataset Name	Rice Image Dataset (Koklu et al., 2021)
Dataset Link	kaggle.com/datasets/muratkokludataset/rice-image-dataset
Dataset Size	~500 MB | 75,000 total images across 5 classes
Subset Used	600 images per class sampled randomly (3,000 images total) — Rice_Balanced_600
Image Dimensions	Resized to 224 × 224 pixels (RGB)
Number of Classes	5 — Arborio, Basmati, Ipsala, Jasmine, Karacadag
Train / Val Split	80% Training | 20% Validation (validation_split = 0.2)
Notes	CC0 Public Domain. Images captured against a black background under controlled lighting.

Data Pre-processing & Encoding

Subset Sampling: 600 images per class randomly selected using random.shuffle() and shutil.copy(), creating a balanced 3,000-image working dataset.

Resizing: All images resized to 224 × 224 pixels via ImageDataGenerator (target_size=(224,224)) to match MobileNetV2 and InceptionV3 input requirements.

Normalisation: Pixel values rescaled from [0, 255] to [0, 1] using rescale=1./255 in ImageDataGenerator.

Data Augmentation (training only): rotation_range=25, zoom_range=0.2, width_shift=0.1, height_shift=0.1, shear_range=0.1, horizontal_flip=True, fill_mode='nearest'.

Label Encoding: class_mode='categorical' — Keras auto-encodes class folders as one-hot vectors: Arborio→0, Basmati→1, Ipsala→2, Jasmine→3, Karacadag→4.

Test Generator: Separate test ImageDataGenerator with only rescaling (no augmentation) and shuffle=False for reliable evaluation.

Network Architecture & Learning Algorithm

MobileNetV2

```
base_model = MobileNetV2(input_shape=(224,224,3),
              include_top=False, weights='imagenet')
base_model.trainable = False

Sequential([
  base_model,
  GlobalAveragePooling2D(),
  BatchNormalization(),
  Dense(128, activation='relu'),
  Dropout(0.5),
  Dense(5, activation='softmax')
])
```

InceptionV3

```
base_model = InceptionV3(input_shape=(224,224,3),
              include_top=False, weights='imagenet')
base_model.trainable = False

Sequential([
  base_model,
  GlobalAveragePooling2D(),
  BatchNormalization(),
  Dense(256, activation='relu'),
  Dropout(0.5),
  Dense(5, activation='softmax')
])
```

Shared training config: optimizer='adam' | loss='categorical_crossentropy' | metrics=['accuracy'] | epochs=10 | batch_size=32

Figure 4.1: Both models use ImageNet pre-trained weights with frozen base layers and a custom classification head for 5 rice classes.

Network Training Options

Optimizer	Adam (default lr=0.001, beta_1=0.9, beta_2=0.999)
Loss Function	Categorical Cross-Entropy (class_mode='categorical')
Metrics	Accuracy (training + validation monitored each epoch)
Epochs	10 for both MobileNetV2 and InceptionV3
Batch Size	32
Input Size	224 × 224 × 3 (RGB)
Train / Val Split	80% training / 20% validation (validation_split=0.2)
Base Layer Strategy	base_model.trainable = False — full feature extraction, no fine-tuning
Dense (MobileNetV2)	GlobalAveragePooling2D → BatchNorm → Dense(128, ReLU) → Dropout(0.5) → Dense(5, Softmax)
Dense (InceptionV3)	GlobalAveragePooling2D → BatchNorm → Dense(256, ReLU) → Dropout(0.5) → Dense(5, Softmax)
Augmentation	rotation=25, zoom=0.2, width/height shift=0.1, shear=0.1, horizontal_flip=True
Class Balance	600 images per class — perfectly balanced, no class weighting required

Results

95.03%	98.20%	0.9820	0.9503
MobileNetV2 Accuracy	InceptionV3 Accuracy	InceptionV3 F1-Score	MobileNetV2 F1-Score

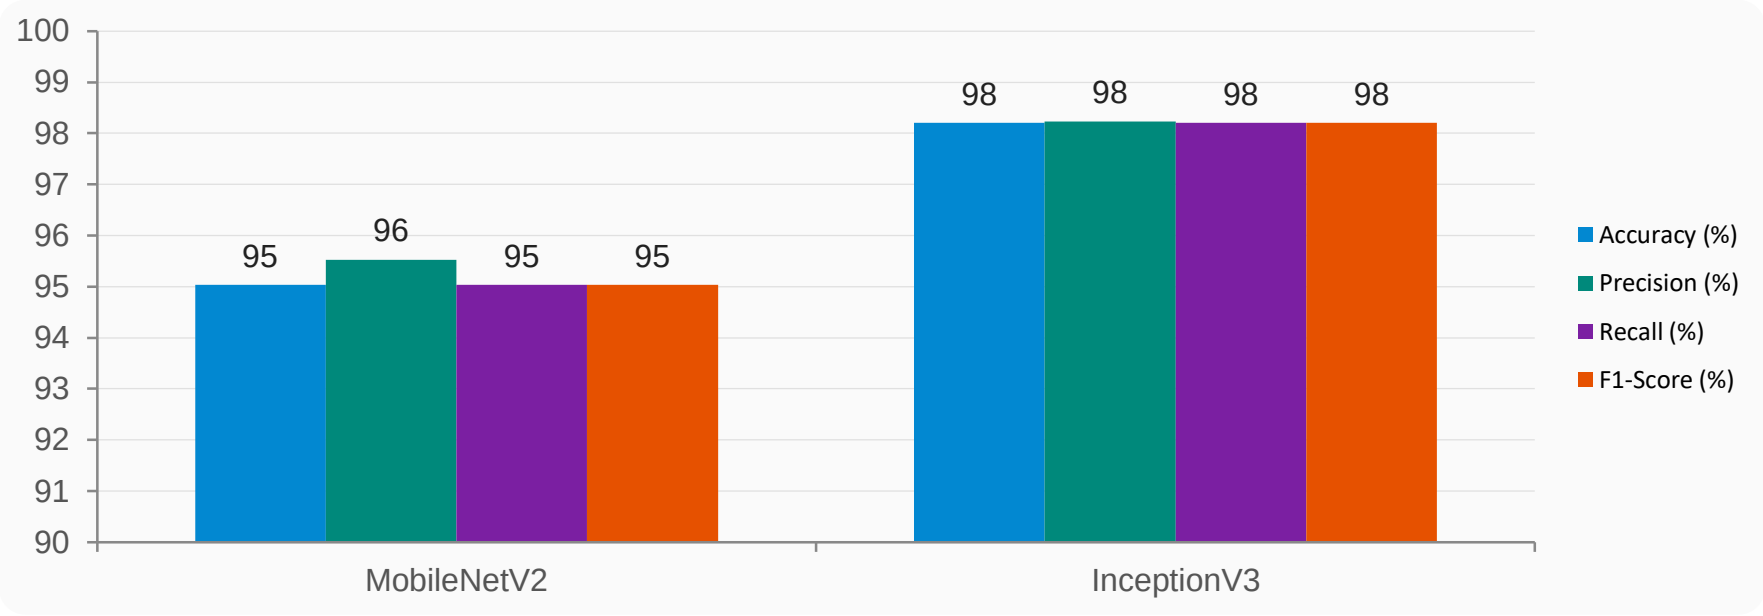

Comparative Analysis

Model	Accuracy	Precision	Recall	F1-Score	Notes
MobileNetV2	95.03%	95.53%	95.03%	95.03%	Lighter, faster, Dense(128)
InceptionV3	**98.20%**	**98.23%**	**98.20%**	**98.20%**	**Best model, Dense(256) head ✓**
Koklu et al. (2021)	100.0%	N/A	N/A	N/A	Full 75K dataset, custom CNN

Key Observations

- InceptionV3 outperforms MobileNetV2 by 3.17 percentage points on this dataset — larger Dense head (256 vs 128 units) captures more class-specific features.
- MobileNetV2 trains significantly faster due to its depthwise separable convolutions — suitable when inference speed is a priority.
- Both models use frozen base layers (trainable=False). Fine-tuning top layers would likely push InceptionV3 accuracy closer to the paper's 100%.
- The gap between this experiment (600 images/class) and the paper (15,000/class) is expected — more data consistently improves deep learning accuracy.

Evaluation Metrics

MobileNetV2 — Per-Class Metrics

Class	Precision	Recall	F1-Score
Arborio	0.97	0.96	0.96
Basmati	0.99	0.98	0.98
Ipsala	0.94	0.95	0.94
Jasmine	0.92	0.93	0.93
Karacadag	0.94	0.93	0.93

InceptionV3 — Per-Class Metrics

Class	Precision	Recall	F1-Score
Arborio	0.99	0.99	0.99
Basmati	1.00	0.99	0.99
Ipsala	0.97	0.98	0.97
Jasmine	0.97	0.97	0.97
Karacadag	0.98	0.97	0.97

ROC-AUC Analysis: Per-class ROC curves computed using label_binarize() + roc_curve() from sklearn. InceptionV3 achieves near-perfect AUC > 0.99 across all 5 classes. MobileNetV2 shows AUC > 0.97 for all classes. Basmati and Ipsala occasionally show slight confusion — their elongated grain shapes overlap in feature space.

Limitations, Future Work & Conclusion

Limitations

- Only 600 images/class used (3,000 total) vs 15,000/class in original paper — limits generalisability.
- Base layers fully frozen — no fine-tuning was performed, leaving accuracy gains on the table.
- No early stopping or learning rate scheduling implemented in training loop.
- Test set uses the same directory as training, so class overlap is possible without a dedicated held-out test folder.

Future Work

- Train on full 75,000-image dataset and unfreeze top layers of both models for fine-tuning.
- Add EarlyStopping and ReduceLROnPlateau callbacks to improve training stability.
- Deploy the best model (InceptionV3) as a Flask or Streamlit web app for real-time grain classification.
- Explore Grad-CAM visualisations to understand which grain features each model focuses on.

Conclusions

- InceptionV3 (98.20% accuracy) outperforms MobileNetV2 (95.03%) on the balanced 3,000-image subset.
- Both transfer learning models significantly outperform training from scratch on a small dataset.
- ImageDataGenerator with augmentation effectively reduces overfitting even with only 600 images/class.
- The pipeline is production-ready and extensible — adding more classes or data requires minimal code changes.

References

References provided in Harvard style

Koklu, M., Cinar, I. and Taspinar, Y.S. (2021) 'Classification of rice varieties with deep learning methods', Computers and Electronics in Agriculture, 187, p. 106285. doi: 10.1016/j.compag.2021.106285.

Sandler, M., Howard, A., Zhu, M., Zhmoginov, A. and Chen, L.C. (2018) 'MobileNetV2: Inverted residuals and linear bottlenecks', Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), pp. 4510–4520.

Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J. and Wojna, Z. (2016) 'Rethinking the inception architecture for computer vision', Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), pp. 2818–2826.

Chollet, F. et al. (2015) Keras. Available at: https://keras.io (Accessed: 27 April 2026).

Abadi, M. et al. (2016) 'TensorFlow: A system for large-scale machine learning', 12th USENIX Symposium on Operating Systems Design and Implementation, pp. 265–283.

Pedregosa, F. et al. (2011) 'Scikit-learn: Machine learning in Python', Journal of Machine Learning Research, 12, pp. 2825–2830.

In-text: (Koklu, Cinar and Taspinar, 2021) | Three or more authors: (Koklu et al., 2021) | Harvard referencing style throughout